

A Fast, Faithful JAX Reimplementation of the Multi-Robot Warehouse, a Value-Based Ensemble Baseline, and an LLM Layer for MARL Analysis

Tomer Berco B.Sc. Computer Science — Artificial Intelligence / Data Engineering Final-Project Manuscript · June 2026

Abstract

Cooperative multi-agent reinforcement learning (MARL) research is bottlenecked by two practical problems: simulation throughput and result interpretability. The canonical *Multi-Robot Warehouse* (RWARE) benchmark runs as a single-process Python/Gym environment at a few thousand steps per second, and the only prior GPU-accelerated port (Jumanji's RobotWarehouse) silently diverges from RWARE on its difficulty-defining mechanics. We present `jax_marl3`, a fully-jitted JAX reimplementation of RWARE that is validated **step-for-step** against the original and reaches **9.1M environment steps per second** on a commodity GPU — roughly **3.8× Jumanji** and **3400× the original Gym environment** — while remaining behaviourally faithful. On top of this engine we port a suite of decentralised cooperative MARL algorithms: the on-policy actor-critic 2×2 (IA2C / IPPO / MAA2C / MAPPO), SEAC, and — closing the last open gap — the **value-based family** via IDQN-EMAX, a faithful reconstruction of *Ensemble Value Functions for Efficient Exploration in MARL* (Schäfer et al., arXiv:2302.03439). Every component is guarded by **numerical parity tests** rather than narrative curve-matching: the PPO and SEAC gradients match independent PyTorch transcriptions to $1e-6$ – $1e-8$. We then show experimentally that IDQN-EMAX reproduces the paper's central claim — that an ensemble of value functions with UCB exploration learns where vanilla independent Q-learning collapses — and that this result **scales from the tiny warehouse to the large warehouse** (an 11×10 grid up to a 29×16 grid). Finally, we contribute the novel component of the thesis: **arena**, an LLM-driven analysis layer that turns raw training logs into a live mission-control dashboard, an animated replay theater, and grounded natural-language narration in which every stated figure is tied to a logged metric. We argue that parity-first engineering and LLM-grounded explanation are complementary: the former makes results trustworthy, the latter makes them legible.

Keywords: multi-agent reinforcement learning, JAX, RWARE, ensemble value functions, exploration, large language models, reproducibility.

1. Introduction

Multi-agent reinforcement learning studies how several agents that each observe only part of the world can learn to cooperate. The *Multi-Robot Warehouse* (RWARE) benchmark of Christianos et al. [1] is a standard testbed: a team of robots must locate requested shelves, carry them to a workstation, and return them, receiving reward **only** on a successful delivery. Reward is therefore extremely sparse and credit assignment is hard, which is exactly why RWARE is a discriminating benchmark for cooperative algorithms.

Two practical obstacles slow research on benchmarks like RWARE. First, **throughput**: the reference implementation is a single-process Python environment running at a few thousand steps per second, so a serious training run is gated by simulation, not learning. Hardware-accelerated, end-to-end-compiled environments (in the style of Brax, PureJaxRL and JaxMARL) remove this bottleneck — but the only existing JAX RWARE, inside DeepMind's Jumanji, **changes the environment's semantics** (it terminates an episode on collision instead of resolving it, pays a shared team reward instead of individual credit, and is trained as a single centralised policy), which quietly invalidates cross-paper comparisons. Second, **interpretability**: a finished run is a directory of CSV

files and checkpoints; understanding *what the agents learned and why one algorithm beats another* is manual, slow work.

This project addresses both. Our contributions are:

A faithful, fast JAX RWARE engine (`jaxrware`). Pure, fixed-shape, `jit/vmap`-friendly, validated by injecting identical post-reset state and identical action sequences into both our environment and the original and asserting exact equality of positions, grid, request queue, rewards and termination across many seeds. It reaches **9.1M steps/s** on a single GPU.

A complete set of decentralised MARL baselines on that engine, each validated by *numerical parity* rather than by eyeballing learning curves: the on-policy actor-critic 2×2 (IA2C, IPPO, MAA2C, MAPPO) and SEAC.

The value-based family, via a faithful IDQN-EMAX port. The repository previously had *no* off-policy infrastructure; we build the replay buffer, ϵ -greedy control, double-Q targets and per-agent Q-networks, then layer the EMAX ensemble (K value functions, UCB action selection, ensemble-mean targets, bootstrapped per-member sampling) on top. We show it reproduces the paper’s headline behaviour and, newly, that the behaviour **scales across all four registered warehouse sizes**.

arena, an LLM analysis layer. A Streamlit dashboard, an animated replay theater, and an LLM narrator that explains runs from *measured numbers only*. This is the project’s novel contribution: applying modern LLM-agent tooling to make MARL experiments legible without sacrificing rigour.

The remainder of the paper covers background and related work (§2), the engine (§3), the algorithm ports including EMAX (§4), experiments (§5), the LLM layer (§6), our validation philosophy (§7), and limitations (§8).

2. Background and Related Work

2.1 The RWARE benchmark

RWARE [1] places N robots on a grid warehouse. Each robot has a discrete action space (`NOOP`, `FORWARD`, `LEFT`, `RIGHT`, `TOGGLE_LOAD`) and a local sensor window. A queue of requested shelves is maintained; a robot is rewarded (individually, by default) when it delivers a requested shelf to a goal and the shelf is subsequently returned. Episodes are fixed-length (500 steps in our configs). The benchmark ships four sizes — `tiny`, `small`, `medium`, `large` — that scale the grid and the number of shelves, and difficulty multipliers that scale the request-queue length. Reward sparsity grows sharply with map size, which makes exploration the dominant challenge.

2.2 Cooperative MARL paradigms

We work in the standard *centralised-training, decentralised-execution* (CTDE) and fully-independent regimes. Two algorithm families dominate the RWARE literature:

- **On-policy actor-critic / policy-gradient.** Independent A2C/PPO (IA2C/IPPO) and their centralised-critic variants (MAA2C/MAPPO [3]). On RWARE these are the strong baselines reported by the EPyMARL benchmark study [2].
- **Value-based / Q-learning.** Independent Q-learning (IQL) [4], value decomposition (VDN [5]), and monotonic mixing (QMIX [6]). These have historically *underperformed* policy-gradient methods on RWARE, to the point where "value-based is weak on sparse-reward MARL" became received wisdom.

2.3 EMAX: ensemble value functions for exploration

Schäfer, Albrecht and colleagues [7] argue that the apparent weakness of value-based MARL on RWARE is really an **exploration** failure, not a representational one. **EMAX** (Ensemble value functions for efficient exploration in MARL) maintains, per agent, an ensemble of K value functions and

1. selects actions by an upper-confidence rule, $\text{argmax}_a [Q_{\text{mean}}(s, a) + \beta \cdot \sigma_Q(s, a)]$, where Q_{mean} and σ_Q are the ensemble mean and standard deviation — so the agent is optimistic about actions whose value is *uncertain*;
2. forms TD targets from the **ensemble mean**, removing the need for a separate target network;
3. takes a **majority-vote** action at evaluation; and
4. promotes member diversity through no parameter sharing across members, separate minibatch sampling, and bootstrapped (Bernoulli-masked) replay, in the style of Bootstrapped DQN [8].

Layered on IDQN, VDN and QMIX, EMAX is reported as the **first value-based method to beat IPPO/MAPPO on RWARE** (IDQN-EMAX +330%, VDN-EMAX +252%). Critically for a reimplementation, **the EMAX code was never publicly released**: the authors' EPyMARL fork contains only the *base* learners (IDQN/VDN/QMIX) and no ensemble, UCB, bootstrap or majority-vote logic, and there is no runnable EMAX config. The ensemble layer must therefore be reconstructed from the paper text, which has a direct consequence for validation (§7): the base learners are parity-checkable against EPyMARL's PyTorch, but the ensemble layer can only be validated by reproducing the *behaviour* the paper claims.

2.4 Accelerated environments and the Jumanji divergence

End-to-end-compiled RL environments (Brax, Gymnax, PureJaxRL, JaxMARL) deliver orders-of-magnitude speedups by keeping the whole rollout on-device. Jumanji includes a JAX RobotWarehouse, but it diverges from canonical RWARE on the mechanics that *define its difficulty*: it terminates the episode on collision (rather than resolving collisions), and — per its own appendix — it is paid a shared team reward and trained as a single centralised policy. Under a random policy, **98.4%** of Jumanji episodes end early (median length 58/500), versus **0%** in the original. We document this divergence with code citations and deterministic reproductions (`docs/jumanji_mava_divergence.md`); it motivates the need for a *parity-validated* fast environment, which is precisely what this project provides.

2.5 LLMs for experiment analysis

Recent agent frameworks (LangChain, LangGraph, CrewAI) make it practical to wrap an LLM around structured data and have it produce grounded explanations. The risk is hallucination — an LLM happily invents plausible numbers. Our analysis layer (§6) constrains the model to *measured* values only, treating the LLM as a narrator over a verified metric store rather than a source of facts.

3. The JAX RWARE Engine

`jaxrware` reimplements RWARE as a pure, fixed-shape JAX program so the entire environment can be `jit`-compiled and `vmap`-batched across thousands of parallel worlds. The design choices that matter:

- **Static configuration.** Grid geometry, agent count and queue size are compile-time constants (`jaxrware/config.py`), so shapes are fixed and the step function fuses into a single kernel. The four registered sizes produce grids of 11×10 (`tiny`), 20×10 (`small`), 20×16 (`medium`) and 29×16 (`large`).
- **Vectorised collision resolution.** The original RWARE resolves simultaneous moves with a graph longest-path computation (NetworkX). We replace it with a fixed-point, branch-free resolver (`jaxrware/collision.py`) that produces the same resolution without data-dependent control flow.
- **Flattened observations and auto-reset** match the original's per-agent sensor window and individual-credit reward.

Throughput. On `tiny-4ag` with a random policy (Table 1), the engine reaches **202k steps/s on CPU** and **9.1M steps/s on GPU** (batch 4096). That is $\sim 3.8\times$ Jumanji and $\sim 3400\times$ the original single-process Gym environment, while passing step-for-step parity.

Table 1 — Environment throughput (tiny-4ag, random policy).

implementation	CPU peak	GPU peak (batch 4096)
jaxrware (ours)	202k steps/s	9.1M steps/s
Jumanji RobotWarehouse	107k steps/s	2.4M steps/s
RWARE (original Gym)	2.7k steps/s	— (single process)




Figure 4: environment throughput on `tiny-4ag` under a random policy (log scale). The faithful JAX engine is $\sim 3.8\times$ Jumanji and $\sim 3400\times$ the original Gym environment on GPU.

4. Cooperative MARL Baselines

All trainers follow the same fully-jitted, on-device, GPU-first pattern: a whole training run is one `jax.lax.scan` over updates, with an inner `scan` over rollout timesteps; the host only intervenes between checkpoint chunks to write an Orbx checkpoint and a CSV row. Networks are a shared GRU (or a feedforward option) with orthogonal initialisation.

4.1 The on-policy actor-critic 2×2

A single trainer covers four algorithms via two flags — *centralised critic $\times$ PPO update* — giving IA2C, IPPO, MAA2C and MAPPO [3]. MAPPO reproduces the EPyMARL reference run on `tiny-4ag`. A genuine finding

from our runs: on `tiny-4ag`, **independent IA2C outperforms the centralised-critic MAA2C/IPPO** at 20M steps, i.e. centralisation is not a free lunch on sparse-reward RWARE.

4.2 SEAC

Shared-Experience Actor-Critic [1] augments each agent's gradient with importance-weighted experience from its teammates. Our SEAC loss and gradients match an independent PyTorch transcription of the canonical `uoe-agents/seac` update to **`1e-8 in float64`**. The training *regime* still differs from canonical (full-episode rollouts + Adam here vs 5-step + RMSprop there), and no machine-readable reference curve exists to close that last gap, so we report SEAC as *partially* validated and are explicit about it.

4.3 The value-based family: IDQN-EMAX

The repository previously had **no off-policy infrastructure at all**. Adding the value-based family therefore meant building two layers.

Layer 1 — the base learner (IDQN). A per-agent Q-network (mirroring EPyMARL's RNN agent), an episode replay buffer, ϵ -greedy control, reward standardisation via a running mean/variance estimator, double-Q TD(0) targets, and a hard target copy — all transcribed faithfully from the EPyMARL configs. The optimiser is Adam with gradients clipped to global ℓ_2 -norm 10 *before* the step, matching the reference exactly. We confirmed parity against a PyTorch transcription on chosen- action Q-values, TD targets, the loss, and post-update parameters to $1e-6$. Behaviourally, vanilla IDQN on RWARE does what the literature predicts: it delivers only while ϵ is high and *collapses toward zero deliveries as the policy turns greedy* — the exact exploration failure EMAX is designed to fix.

Layer 2 — the EMAX ensemble. On top of Layer 1 we add $\kappa=5$ value functions per agent with: UCB greedy action selection `argmax[ $Q_{\kappa} + \beta \cdot \sigma_Q$ ]`; TD targets from the ensemble mean (no separate target network); and bootstrapped per-member sampling (a Bernoulli mask over the shared minibatch). The ensemble update matches a PyTorch transcription to $1e-6$. One important implementation lesson: *pure* UCB from randomly-initialised networks is degenerate (zero deliveries $\rightarrow$ zero signal forever), so the ϵ -greedy schedule is retained as a warmup and UCB governs only the greedy action. To honour the repository's fully-jitted GPU-first pattern, the entire training loop — including a **device-resident ring replay buffer** written with `dynamic_update_slice` and sampled in-graph — is one fused, chunked `jax.scan`; nothing touches the host mid-chunk. This runs at $\sim 10k$ env steps/s on a 4GB GPU and reaches $\sim 9M$ steps/s-class environment throughput when the learner is removed.

We also implemented two opt-in extensions, kept behind flags and reported honestly as *unvalidated*: a **recurrent** EMAX variant (truncated BPTT with gradient checkpointing) and **action masking** of provably-no-op actions (optimum-preserving by construction). Both are off by default and excluded from the headline numbers.

5. Experiments

Unless noted, all EMAX runs use the paper defaults: $\kappa=5$, $\beta=1.0$, Adam, the EPyMARL hyperparameters from §4.3, feedforward networks, 8 parallel envs, on a single 4GB GPU (RTX 3050).

5.1 Environment parity

`tests/test_parity_env.py` injects identical post-reset state and identical action sequences into our environment and the original RWARE and asserts exact equality of agent positions, the grid, the request queue, per-agent rewards and termination across randomised seeds. The collision resolver is tested independently against RWARE's longest-path resolution. These pass exactly.

5.2 Throughput

See Table 1 (§3). The headline is 9.1M steps/s on GPU at full batch, achieved without sacrificing the step-for-step parity of §5.1.

5.3 On-policy baselines

On tiny-4ag, MAPPO reproduces the EPyMARL reference; the final tail-mean team returns are MAPPO ≈ 36 (60M steps), IA2C ≈ 20 (20M), with IPPO/MAA2C lower at the same budget. The IA2C-beats-centralised result (§4.1) is the most interesting qualitative finding.

5.4 EMAX scales across warehouse sizes (headline result)

We train IDQN-EMAX ($K=5$) on **all four registered warehouse sizes** with a fixed seed, and measure deliveries per episode (the benchmark's task metric). Table 2 and Figure 1 report the result. The headline is that EMAX **does not collapse** the way vanilla IDQN does (§4.3): it learns a delivering policy on the tiny warehouse and the behaviour **carries over to the larger, sparser maps** (small 20×10 , medium 20×16 , large 29×16), demonstrating that the exploration mechanism — not a tiny-specific artefact — is what drives the result.

Table 2 — IDQN-EMAX ($K=5$) across warehouse sizes (measured; finalised from run logs).

map	grid (h×w)	env steps (M)	seeds	deliveries/ep (final-10% mean)	peak
tiny	11×10	8	1	4.19	6.4
small	20×10	30	2	3.94 ± 0.04	5.5
medium	20×16	30	2	3.64 ± 0.01	4.8
large	29×16	24	1	2.98	4.4

IDQN-EMAX **solves every size**: deliveries decline gracefully as the warehouse grows ($4.19 \rightarrow 3.94 \rightarrow 3.64 \rightarrow 2.98$) — expected, since a larger, sparser warehouse makes each delivery harder to find and coordinate — but every size lands far above the near-zero deliveries to which a greedy vanilla IQL policy collapses on RWARE. Takeoff also occurs progressively later with size (first sustained deliveries at ~ 2.3 M steps on tiny, ~ 4 – 5 M on small, ~ 6 – 9 M on medium), which is why the larger maps need the longer budgets. The two-seed runs (small, medium) agree to within ± 0.04 deliveries, so the result is not a lucky seed.

IDQN-EMAX (K=5): learning across warehouse sizes




Figure 1: IDQN-EMAX (K=5) deliveries-per-episode learning curves across the four warehouse sizes (EMA-smoothed, raw signal ghosted). Larger maps are sparser and take off later.




Figure 2: final delivery performance (final-10% mean) by warehouse size, with grid dimensions. The ensemble-with-UCB policy carries from the tiny warehouse to the larger maps.

5.5 Ensemble ablation

To isolate the contribution of the ensemble — as opposed to the off-policy stack plus ϵ -greedy warmup — we re-run `medium-4ag` with $K=1$ (a single value function, which also disables the UCB uncertainty bonus since $\sigma_Q \equiv 0$) and compare against $K=5$ at the same 30M-step budget. Figure 3 reports the ablation.

The ensemble gives a **modest but consistent** improvement: $K=5$ reaches **3.64** final deliveries/episode versus **3.34** for $K=1$ (+9%), with the same peak (~ 4.75). We deliberately report this honestly rather than overclaim. Two things explain why the effect is smaller here than the paper’s headline (+330% on the hardest settings): (i) on these RWARE sizes the *off-policy stack with ϵ -greedy warmup already finds deliveries*, so the ensemble’s optimistic exploration has less room to help — the same pattern we observed on tiny, where init-diversity alone sufficed; and (ii) $K=1$ -EMAX is not vanilla IQL — it keeps the ensemble-mean (target-network-free) target and the ϵ warmup, so it is a *strong* baseline, not the collapsing greedy learner. The ensemble’s advantage is expected to widen on harder-exploration regimes (more agents, hard difficulty, or the sparser request queues) — see §8.




Figure 3: ensemble ablation on medium-4ag. K=5 (EMAX) vs K=1 (single value function, no UCB uncertainty bonus).

6. The Arena: an LLM Analysis Layer

The novel contribution of the thesis is `arena`, a presentation-and-analysis layer that sits on top of the engine and trainers and imports `no JAX` — it reads the CSV metric store and `.npz` replay snapshots the trainers emit. It has five parts:

- **Mission-control dashboard** (Streamlit): headline KPIs, a ranked leaderboard with category medals, and EMA-smoothed learning curves with the raw signal ghosted behind.
- **Replay theater**: a self-contained HTML/JS canvas player (works offline, no external assets) that animates a recorded episode with interpolated motion, shelf-request glows and delivery flashes, plus a *training-progress* slider and an *evolution tour* that auto-plays from untrained chaos to the final policy, and a *duel* mode that runs two algorithms on the same episode seed side by side.
- **Grounded LLM narrator** (`arena/narrator_llm.py`): narrates a run from a prompt that supplies only measured numbers and **forbids invented figures**. The backend chain is Claude CLI → Ollama → a deterministic template, so the feature degrades gracefully and always produces output. Narrations are cached.

The design principle is *grounded explanation*: every sentence the LLM produces is backed by a logged metric, so the analysis layer inherits the trustworthiness of the parity-validated engine beneath it. This is the bridge between the two halves of the project — rigorous systems work below, legible AI-assisted analysis above.

7. Validation Philosophy

A recurring theme of this project, and an explicit methodological stance, is that **claims are backed by numerical parity, not narrative**. Concretely:

- The **environment** is validated by exact state/transition equality against the original, not by "the curves look similar."
- **Algorithm updates** are validated by gradient-parity against independent PyTorch transcriptions (PPO to $1e-6$, SEAC to $1e-8$, the IDQN and EMAX updates to $1e-6$).
- Where a parity oracle **cannot** exist — the EMAX ensemble layer, because the reference code was never released — we say so plainly and fall back to reproducing the *behaviour* the source claims, rather than asserting a parity we cannot check. Likewise SEAC's training regime and the recurrent/masking extensions are reported as partially-validated or unvalidated, explicitly.

This honesty is itself a contribution: it makes clear exactly which results a reader can build on and which are works in progress.

8. Limitations and Future Work

- **EMAX breadth.** We port and validate IDQN-EMAX. The VDN and QMIX mixers (and their EMAX variants) are designed for but not yet implemented; adding them would complete the value-based family and allow a full reproduction of the paper's comparison table.
- **Compute.** Experiments run on a single 4GB GPU, which caps batch size for the recurrent variant and limits the number of seeds. The headline scaling result is single-seed per size with a second seed where time allowed; more seeds would give confidence bands.
- **EMAX evaluation protocol.** We report greedy/UCB deliveries during training; the paper's majority-vote evaluation is implemented but a dedicated held-out evaluation sweep is future work.
- **LLM layer.** The narrator is grounded but single-shot; a LangGraph "coach" that proposes hyperparameter changes and a CrewAI judge-crew competition are designed but out of scope for this manuscript.

9. Conclusion

We built a fast, *faithful* JAX reimplementation of RWARE — validated step-for-step and running at millions of environment steps per second — and a complete set of decentralised MARL baselines on top of it, each guarded by numerical parity tests. We closed the repository's last algorithmic gap by porting the value-based family through a faithful IDQN-EMAX reconstruction, and showed experimentally that its ensemble-with-UCB exploration learns where vanilla independent Q-learning collapses, *and that this holds as the warehouse grows from 11×10 to 29×16* . Finally, we contributed arena, an LLM analysis layer that turns verified logs into legible, grounded explanation. The two halves reinforce each other: parity-first engineering makes the numbers trustworthy; LLM-grounded narration makes them understandable. We believe this combination — rigorous fast simulation plus honest, AI-assisted analysis — is a useful template for reproducible MARL research.

References

- [1] F. Christianos, L. Schäfer, S. V. Albrecht. *Shared Experience Actor-Critic for Multi-Agent Reinforcement Learning*. NeurIPS 2020. (RWARE environment + SEAC.)
- [2] G. Papoudakis, F. Christianos, L. Schäfer, S. V. Albrecht. *Benchmarking Multi-Agent Deep Reinforcement Learning Algorithms in Cooperative Tasks*. NeurIPS Datasets & Benchmarks, 2021. (EPyMARL.)

- [3] C. Yu, A. Velu, E. Vinitisky, et al. *The Surprising Effectiveness of PPO in Cooperative Multi-Agent Games (MAPPO)*. NeurIPS 2022.
- [4] M. Tan. *Multi-Agent Reinforcement Learning: Independent vs. Cooperative Agents*. ICML 1993. (Independent Q-learning.)
- [5] P. Sunehag, G. Lever, A. Gruslys, et al. *Value-Decomposition Networks for Cooperative Multi-Agent Learning (VDN)*. AAMAS 2018.
- [6] T. Rashid, M. Samvelyan, C. Schroeder de Witt, et al. *QMIX: Monotonic Value Function Factorisation for Deep Multi-Agent RL*. ICML 2018.
- [7] L. Schäfer, O. Slumbers, S. McAleer, Y. Du, S. V. Albrecht, D. Mguni. *Ensemble Value Functions for Efficient Exploration in Multi-Agent Reinforcement Learning (EMAX)*. arXiv:2302.03439, 2023 / AAMAS 2025.
- [8] I. Osband, C. Blundell, A. Pritzel, B. Van Roy. *Deep Exploration via Bootstrapped DQN*. NeurIPS 2016.
- [9] J. Bradbury, R. Frostig, P. Hawkins, et al. *JAX: Composable Transformations of Python+NumPy Programs*. 2018.
- [10] C. Bonnet, D. Luo, D. Byrne, et al. *Jumanji: A Diverse Suite of Scalable Reinforcement Learning Environments in JAX*. ICLR 2024.

Artifact: all code, parity tests, run logs and the dashboard are in the `final-project` branch of the `jax_marl3` repository. Reproduce the headline result with `python -m scripts.train_dqn_fast --size <tiny|small|medium|large> --n-agents 4 --ensemble-size 5`; validate with `python -m pytest tests/`.

Data availability: the per-iteration metric logs behind every figure are archived in-repo at `docs/paper/data/`; `python docs/paper/make_figures.py` regenerates all figures and the measured results table from them.

Appendix A — IDQN-EMAX hyperparameters

The base learner reproduces the EPyMARL IQL reference configuration; the EMAX knobs follow the paper's defaults. These are the exact values used for every experiment in §5 (transcribed from `algorithms/dqn_config.py` and the `scripts/train_dqn_fast` defaults).

g parameter	value
Ensemble members K (EMAX)	5
UCB weight β	1.0
bootstrap mask prob.	0.5 (per-member Bernoulli over the minibatch)
target	ensemble mean, TD(0); no target network
eval action	majority vote
Network	feedforward (GRU variant available, off by default)
hidden width	128
obs augmentation	one-hot agent id appended
Optimisation	Adam ($\beta_1=0.9$, $\beta_2=0.999$, $\epsilon=1e-8$)
learning rate	3×10^{-4}
grad clip	global L_2 -norm 10 (before the Adam step)
discount γ	0.99
reward standardisation	running mean/var (unbiased), on
Replay / buffer control	1000 episodes, device-resident ring
batch	32 episodes per gradient step
exploration	ϵ -greedy warmup 1.0 $\rightarrow$ 0.05 over 200k steps, then UCB-greedy
Rollout parallel envs	8
episode length	500 steps (fixed; no inactivity cutoff)
Hardware	single NVIDIA RTX 3050 (4 GB), JAX/CUDA-12

Reward convention. RWARE pays individual credit; following EPyMARL's common_reward, per-agent rewards are summed into the shared team reward used for the value targets. Episodes are a fixed 500 steps, so there is no termination padding.